

Politechnika Warszawska

W Y D Z I A Ł   M A T E M A T Y K I  
I   N A U K   I N F O R M A C Y J N Y C H



# Praca dyplomowa magisterska

na kierunku Informatyka

Rejection Option in Pattern Recognition Problem - Selected Issues

inż. Piotr Waszkiewicz

Numer albumu 254218

promotor

dr hab. inż. Władysław Homenda

WARSZAWA 2017

.....

podpis promotora

.....

podpis autora

## Abstract

An analysis of the presented study seeks a solution to a common problem in a classification issue, which is detecting and rejecting data not suited for classification. Contaminated data that emerges from noisy environment can lead to a situation in which even well trained models yield bad results. This is a serious problem for processes that rely on a classifiers' efficiency in which rejecting received data is more acceptable than classifying it wrongly, e.g. tumour detection algorithm should refuse to make medical evaluation of provided image if it is too blurry rather than trying to guess patient's health condition.

Although artificial intelligence gained much importance and is used in many aspects of humans life (even outside of pure scientific fields), there's still a need for newer approaches and methods. Commonly used algorithms and models change very frequently as new problems arise. Study presented in this thesis introduces modifications to some of the oldest and well known techniques and tries to combine them in order to create tools with much higher capabilities.

**Keywords:** classification, rejection, svm, random forest, knn, ellipsoid, artificial intelligence



## Streszczenie

Jednym z problemów z dziedziny sztucznej inteligencji jest rozpoznawanie i klasyfikacja elementów dotyczących rozwiązywanego problemu. Określenie przynależności nieznanego wzorca jest zagadnieniem o tyle problematycznym, że wymaga ono zadbania o dobrej jakości dane. Odczyty zaszumione, nienależące do dziedziny badanego problemu mogą skutkować błędnymi wynikami. Taka niepoprawna klasyfikacja może mieć poważne konsekwencje - stwierdzenie braku zmian rakowych w organizmie pacjenta z powodu źle wykonanego badania pociąga za sobą o wiele większe konsekwencje niż brak jakiegokolwiek diagnozy.

Celem niniejszej pracy jest przeprowadzenie badań i próba skonstruowania nowych rodzajów klasyfikatorów które posiadać będą cechę odrzucania elementów obcych, czyli nienależących do rozwiązywanego problemu. Chociaż sztuczna inteligencja jest obecnie prężnie rozwijającą się dziedziną w której nowe podejścia i algorytmy pojawiają się bardzo często, nie można zapominać jak duże znaczenie wciąż mają stare rozwiązania. Podczas przeprowadzanych testów na potrzeby tej pracy wykorzystywane będą właśnie jedne z najstarszych, a jednocześnie najbardziej popularnych, klasyfikatorów t.j. maszyna wektorów wspierających, lasy losowe czy algorytm knn.

**Słowa kluczowe:** klasyfikacja, odrzucanie, svm, lasy losowe, knn, elipsoida, sztuczna inteligencja



inż. Piotr Waszkiewicz

Warsaw, .....

Nr albumu 254218

### Declaration

I hereby declare that the thesis entitled „Rejection Option in Pattern Recognition Problem - Selected Issues”, submitted for the magisters degree, supervised by dr hab. inż. Władysław Homenda, is entirely my original work apart from the recognized reference.

.....

inż. Piotr Waszkiewicz





# Contents

<b>List of Figures</b>	
<b>List of Tables</b>	
<b>1. Introduction</b>	<b>1</b>
<b>2. Common Classifiers</b>	<b>2</b>
2.1. Implementation	2
2.2. kNN	3
2.3. SVM	4
2.4. Random Forest	6
<b>3. Quality Evaluation</b>	<b>9</b>
<b>4. Datasets</b>	<b>10</b>
<b>5. Classifier Trees</b>	<b>11</b>
5.1. Balanced Tree	11
5.1.1. Structure	11
5.1.2. Classifiers creation	12
5.1.3. Classification rules	13
5.1.4. Implementation details	13
5.2. Slanting Tree	14
5.2.1. Structure	14
5.2.2. Classifiers creation	14
5.2.3. Classification rules	15
5.2.4. Implementation details	15
5.3. Slanting Tree with ordered classes	16
5.3.1. Description	16
5.3.2. Implementation details	16
5.4. Slanting Tree 2	18
5.4.1. Description	18
5.4.2. Implementation details	18

5.5. Results . . . . .	18
5.6. Summary . . . . .	21
<b>6. Classifier Arrays . . . . .</b>	<b>22</b>
6.1. One-versus-all . . . . .	22
6.1.1. Description . . . . .	22
6.1.2. Implementation details . . . . .	23
6.2. One-versus-one . . . . .	23
6.2.1. Description . . . . .	23
6.2.2. Implementation details . . . . .	24
6.3. One-versus-one modified . . . . .	24
6.4. Results . . . . .	24
<b>7. Identifier Arrays . . . . .</b>	<b>26</b>
7.1. Minimum Volume Enclosing Ellipsoid . . . . .	26
7.2. Ellipsoids array classifier . . . . .	27
7.3. Results . . . . .	28
 <b>A. An Appendix . . . . .</b>	
 <b>Bibliography . . . . .</b>	

## List of Figures

2.1	Visualization of area coverage of three different class membership for kNN classifier with $k=15$ , using euclidean metric. Image taken from [1]	3
2.2	SVM hyperplane construction with the biggest possible margin for training dataset. Image taken from [1]	5
2.3	Different class area coverages resulting from usage of different kernel functions. Image taken from [1]	5
2.4	A funny example of a decision tree	6
2.5	Visualization of a random forest consisting of $B$ different decision trees	7
4.1	Visualization of scanned digits from MNIST database, image taken from [2]	10
5.1	Balanced Tree structure obtained during real life tests	12
5.2	Balanced Tree rejection scheme in leaf node	13
5.3	Bottom part of the Slanting Tree with nodes for classes 8 and 9 (nodes for classes from 0 - 7 not visible).	17
6.1	“one-versus-all” rejection method. Unknown pattern passes through an array of specially prepared classifiers, one for each class. If each classifier says it is not native, it is rejected.	22
6.2	“one-versus-all” rejection method. Unknown pattern passes through an array of specially prepared classifiers, one for each class. If each classifier says it is not native, it is rejected.	23
7.1	Array of ellipsoids that works as a classifier with rejection capabilities. Elements outside of existing ellipsoids are treated as a foreign patterns.	28



## List of Tables

3.1	Quality measures for classification with rejection. . . . .	9
5.1	Example empty result matrix . . . . .	19
5.2	Measures values (described in Chapter 3) for classifier trees using various common classifiers on training data (described in Chapter 4) . . . . .	20
5.3	Measures values (described in Chapter 3) for classifier trees using various common classifiers on test data (described in Chapter 4) . . . . .	20
6.1	Measures values (described in Chapter 3) for classifier hierarchy arrays using various common classifiers on training data (described in Chapter 4) . . . . .	25
6.2	Measures values (described in Chapter 3) for classifier hierarchy arrays using various common classifiers on test data (described in Chapter 4) . . . . .	25
7.1	Results obtained for Ellipsoid arrays . . . . .	28



# 1. Introduction

Study presented in this paper tries to combine few selected classifiers in such way that will empower them to gain rejection capabilities. The main goal is to come up with a model which structure allows to reject patterns that are outside of native elements classes' scope without any prior knowledge about such patterns. Those outliers, denoted as foreign elements, are very common in real life situations when dealing with noisy, erroneous or unknown measurements. The main drawback of commonly used classifiers is their inability to reject foreign elements without any knowledge about them. Classifiers such as SVM, random forest or kNN (described better in Chapter 2) must always put presented pattern to one of the classes they were trained on. This requirement to always classify provided pattern to one of the classes their were trained on forces inclusion of foreign elements within training sets if there's a need to reject certain elements. Although not impossible, this approach is quite impractical as most of the time there are just too many possible cases of foreign elements. This paper tries to find solution to this presented problem.

The document is structured as follows: the first chapter introduces reader to the commonly used, well-known classifiers, describing their structures and briefly explaining the way they work. The next chapter covers quality measurements which help in evaluating proposed solutions and their results. Next, the datasets used during testing process are described. The following chapters present proposed classifier structures with rejection option, explaining the main ideas behind them and comparing their results achieved on provided data.

## 2. Common Classifiers

The task of classification aims at categorising unknown elements to their appropriate groups. The procedure is based on quantifiable characteristics obtained from the source signal. Those characteristics, i.e. features, are gathered in a feature vector (a vector of independent variables) and each pattern is described with one feature vector. It is expected that patterns accounted to the same category are in a relationship with one another. In other words, subjects and objects of knowledge accounted to the same category are expected to be in some sense similar. There are many mathematical models that can be used as classifiers, such as SVM, random forest, kNN, regression models, or Neural Networks. Their main disadvantage lies in their need to be trained prior to usage, which makes them unable to recognize elements from a new class, not present during the training process. This behaviour can be especially troublesome in an unstable, noisy environment, where patterns sent for classification can be corrupted, distorted or otherwise indistinguishable.

### 2.1. Implementation

Implementations of the common classifiers described in this chapter were taken from scikit-learn<sup>1</sup> Python library[3]. It is a popular, open source project using BSD license and built on NumPy<sup>2</sup>, SciPy<sup>3</sup> and matplotlib libraries. The project was started in 2007 by David Cournapeau as a Google Summer of Code project and is currently maintained by a team of volunteers. The library contains implementations of many algorithms to be used, among others, in classification, regression, clustering, dimensionality reduction and preprocessing problems.

---

<sup>1</sup>scikit-learn webpage: <http://scikit-learn.org>

<sup>2</sup>NumPy webpage: <http://www.numpy.org>

<sup>3</sup>SciPy webpage: <https://www.scipy.org>



## 2.2. kNN

The k-Nearest Neighbours algorithm, denoted as kNN, is an example of a “lazy classifier”, where the entire training dataset is the model. There is no typical model building phase, hence the name. Class membership is determined based on class labels encountered in  $k$  closest observations in the training dataset, [4]. In a typical application, the only choice that the model designer has to make is selection of  $k$  and distance metrics. Both are often determined experimentally with a help of supervised learning procedures. Example of area coverage for three classes used in kNN classification issue can be seen in Figure 2.1.

The kNN classifier implementation available within scikit-learn package allows to make adjustments to certain parameters that are crucial in classification issue:

- *n\_neighbors* - corresponds to the  $k$  value, determines number of nearest points used to classify pattern
- *metric* - the distance metric to use for the tree




Figure 2.1: Visualization of area coverage of three different class membership for kNN classifier with  $k=15$ , using euclidean metric. Image taken from [1]

### 2.3. SVM

Support Vector Machines (SVM) are a collection of supervised learning methods used for classification, regression and outliers detection. The SVM algorithm relies on a construction of hyperplane with a maximal margin that separates patterns of two classes [5]. Creation of the hyper-plane that has the largest distance to the nearest training data points of any class (so-called functional margin) is important since, in general, the larger the margin the lower the generalization error of the classifier.

In SVM's mathematical definition the two classes' labels are denoted as -1 and 1. When treating elements from those sets as points of the Euclidean space  $\mathbb{R}^n$  (or vectors of this space) the SVM training can be seen as the problem of finding the maximum-margin hyperplane that divides those samples. This issue can be described by formula:

$$w * x - b = 0$$

where  $w, x \in \mathbb{R}^n, b \in \mathbb{R}$ . The  $x_i$  vectors are samples from the training set, and  $w$  is a normal vector to the hyperplane, obtained as a linear combination of those training vectors that lie at borders of the margin:

$$w = \sum_i \alpha_i x_i$$

Those of the training vectors  $x_i$  that satisfy the following condition:

$$y_i(x * x_i - b) = 1$$

are called support vectors, and have their corresponding  $\alpha_i \neq 0$ . The  $y_i \in -1, 1$  corresponds to the class labels that training data consists of. The linear decision function used for classifying patterns is expressed as follows:

$$I(x) = \text{sgn}(\sum \alpha_i x_i * x - b)$$

where  $\alpha_i x_i = w_i$ . SVM efficiency can be enhanced by using different kernel functions which help in solving non-linearly-separable problems. The generalized decision function using kernel function  $K$ :

$$I(x) = \text{sgn}(\sum \alpha_i K(x_i, x) - b)$$

SVMs are effective in high-dimensional spaces, memory efficient, and quite versatile because of the many kernel functions that can be specified for the decision function. Implementation available as part of scikit-learn package lets user specify and tweak many aspects of classifier such as:




Figure 2.2: SVM hyperplane construction with the biggest possible margin for training dataset. Image taken from [1]




Figure 2.3: Different class area coverages resulting from usage of different kernel functions. Image taken from [1]

- $C$  - penalty parameter  $C$  of the error term, used to regularize the estimation. If dealing with noisy observations it's recommended to decrease its value
- *kernel* - kernel type used in the algorithm, in this paper one of "poly" or "rbf" values are used. "poly" stands for polynomial kernel using following equation  $(\gamma \langle x, x' \rangle + r)^d$  (where  $d$  is function degree, with default value 3), "rbf" is an acronym for radial basis function with given equation  $\exp(-\gamma |x - x'|^2)$
- *gamma* - kernel coefficient for "rbf", "poly" types as can be seen in the kernel equations
- *degree* - degree of the polynomial kernel function

It is worth noting though that in some cases, where the number of features is much greater than the number of samples, using support vector machines can give poor results, and is not cost-efficient when calculating probability estimates.

## 2.4. Random Forest

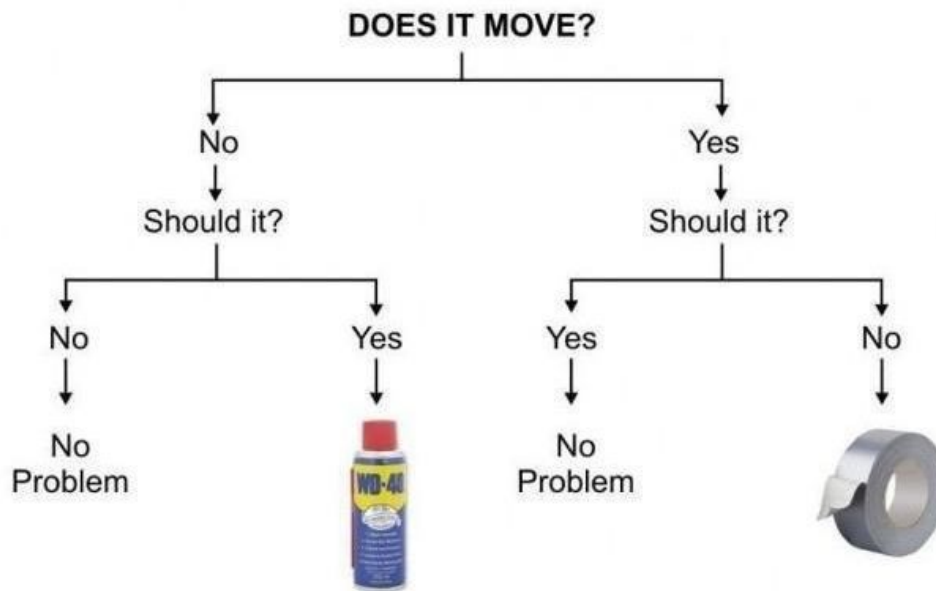


Figure 2.4: A funny example of a decision tree

Random forest is a popular ensemble method. The main principle behind ensemble methods, in general, is that a group of “weak learners” can come together to form a “strong learner”. In the random forest algorithm [6] the weak learners are decision trees, which are used to predict class labels. A decision tree is a decision support tool that uses a tree-like graph for classification issue.

Each graph node performs a test on an attribute of the provided pattern and sends it to its child node via a branch that represents the outcome of the test. Each leaf in a decision tree represents a certain class label. In other words for a feature vector representing one pattern a decision tree calculates its class label by dividing value space into two or more subspaces. More precisely, an input data is entered at the top of the tree and as it traverses down the tree the data gets bucketed into smaller subsets. There are many advantages of using decision trees. Their results are easy to interpret and visualize in form of a graph, they can handle multi class classification problems and perform well even if its assumptions are somewhat violated by the true model from which the data were generated. On the other hand, the main drawbacks connected to their usage consist of overfitting problem caused by creating too complex trees on a very complicated data, and instability caused by small variations in the data that might result in a completely different tree being generated. That last problem is easily mitigated by ensembling set of decision trees into a random forest.




Figure 2.5: Visualization of a random forest consisting of  $B$  different decision trees

In the random forest a large number of classification trees is formed, which altogether serve as a classifier. In order to grow each tree, a random selection of rows from the training set is drawn. Random sampling with replacement is also called bootstrap sampling. In addition, when constructing trees for a random forest at each node  $m$  variables out of the set of all input variables are randomly selected, and the best split on these  $m$  is used to split the node. After a relatively large number of trees is generated, they vote for the most popular class. Some of the parameters used for improving classification rates that are available within scikit-learn package random forest implementation:

- *n\_estimators* - determines number of trees used by random forest in the algorithm
- *max\_depth* - the maximum depth of each tree in the forest
- *max\_features* - the number of features to consider when looking for the best split
- *min\_samples\_leaf* - the minimum number of samples required to be at a leaf node

Random forests join few important benefits: (a) they are relatively prone to the influence of outliers, (b) they have an embedded ability of feature selection, (c) they are prone to missing values, and (d) they are prone to over-fitting.

### 3. Quality Evaluation

In order to evaluate the quality of the proposed methods a set of measures is used, described below and in Table 3.1.

- *Correctly Classified* is the number of native patterns classified as native with a correct class label.
- *True Positives* is the number of native patterns classified as native (no matter, into which native class).
- *False Negatives* is the number of native patterns incorrectly classified as foreign.
- *False Positives* is the number of foreign patterns incorrectly classified as native.
- *True Negatives* is the number of foreign patterns correctly classified as foreign.

Table 3.1: Quality measures for classification with rejection.

Native Precision	=	$\frac{TP}{TP+FP}$	Accuracy	=	$\frac{TP+TN}{TP+FN+FP+TN}$
Foreign Precision	=	$\frac{TN}{TN+FN}$	Strict Accuracy	=	$\frac{CC+TN}{TP+FN+FP+TN}$
Native Sensitivity	=	$\frac{TP}{TP+FN}$	Fine Accuracy	=	$\frac{CC}{TP}$
Foreign Sensitivity	=	$\frac{TN}{TN+FP}$	Strict Native Sensitivity	=	$\frac{CC}{TP+FN}$
F-measure = $2 \cdot \frac{\text{Precision} \cdot \text{Sensitivity}}{\text{Precision} + \text{Sensitivity}}$					

## 4. Datasets

All quality measures described in Chapter 3 obtained and presented in this paper were calculated from classifiers' results for certain datasets. Those sets, referred to as native and foreign, are the result of applying feature-extraction function to images containing digits and letters. The original data comes from the well-known MNIST database[7].



Figure 4.1: Visualization of scanned digits from MNIST database, image taken from [2]

The native set consists of 10,000 scanned digit images, with ten different classes one for every digit (0 - 9) and approximately 1000 samples for each class. This set is further divided into training and test sets in 7:3 ratio. The foreign set consists of 26,000 images of scanned letters and is not divided internally because it is used only in rejection option evaluation.

Every pattern within those two datasets consists of 24 unique features that were extracted to ensure best classification capabilities. Examples of features are: maximum/position of maximum values of projections, histograms of projections, transitions, offsets; raw moments, central moments, Euler numbers etc.



## 5. Classifier Trees

Common classifiers described in the Chapter 2 return results in form of a class label that provided pattern was classified to. Such approach leaves no room for estimating class-belonging probabilities which, in return, results in inability to reject provided data, treating it as an outlier. By combining those classifiers and organising them in a complex structures it is possible to create objects with unique rejection capabilities in exchange for slightly increased pattern-processing time. This chapter describes such structures, shaped in form of binary trees.

### 5.1. Balanced Tree

#### 5.1.1. Structure

The main idea behind Balanced Tree structure is to create a graph tree in which every path from root to leaf consists of increasingly precise classifiers. What it means is that every pattern, that should be classified, is tested against certain number of common classifiers, where each subsequent one is clarifying this unknown pattern's affiliation to one of the classes.

The Balanced Tree construction begins with creation of a root node which represents a situation in a classification process in which all possible class memberships for an unknown pattern are taken into account. It can be said that the root of the Balanced Tree represents a set consisting of all classes in the training set, because it is yet unknown from which class a pattern would be. The process of clarifying pattern's class belonging starts by designating the central points for each class in the set of classes represented by this node. This is done by calculating arithmetic average of all points from certain class set:

$$p_{central} = \frac{\sum_{i=1}^n p_i}{n}$$

where  $n$  is number of elements  $p_i$  belonging to certain class in the dataset. Next step involves using clustering algorithm to divide all of those central points into two distinctive sets. The idea is to group those class representatives that are most similar to each other. The process of Balanced Tree structure creation is continued further by passing two classes sets designated by clustering

algorithm, one to each child nodes. The process of new node creation is then applied to each of those two child nodes and continued until there is only one class left. A node representing only one class cannot use clustering method because there is insufficient number of classes to divide, and so it becomes the tree leaf.

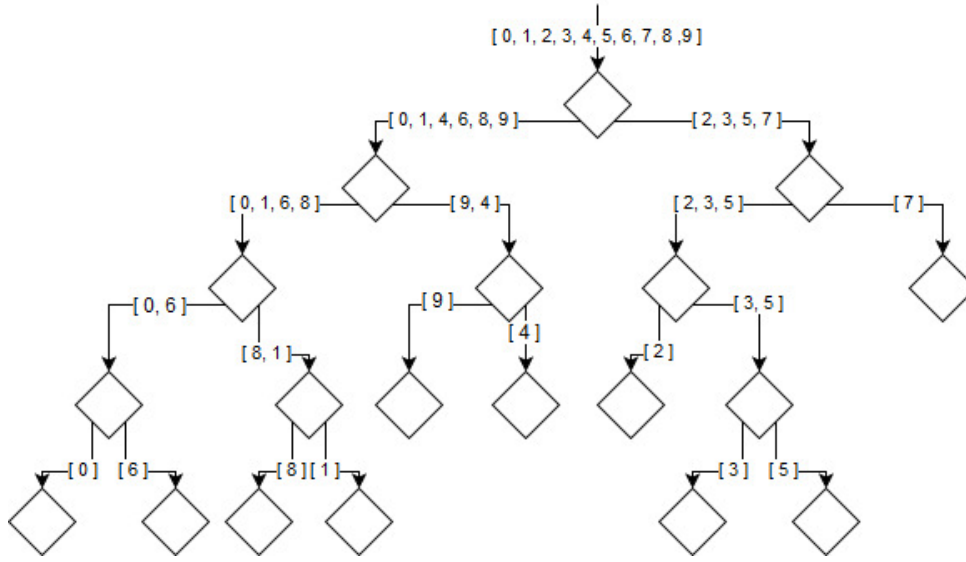


Figure 5.1: Balanced Tree structure obtained during real life tests

### 5.1.2. Classifiers creation

After finishing Balanced Tree architecture creation each non-leaf node is assigned a binary classifier trained on a data consisting of a training samples from classes assigned to this particular node. Those classes that are represented by its left child node are joined together and treated as a class '0', while the ones in the right child node are labelled as a class '1'. For example, if there are four classes assigned to a certain tree node, labelled as  $a, b, c, d$ , and a clustering method divided them into two sets  $a, d$  (assigned to left child node) and  $b, c$  (assigned to right child node), the classifier will be trained on data samples treating points from classes  $a, d$  as if they all were from an artificial class '0' and classes  $b, c$  as class '1'. The only issue that arises from such attitude is inability for leaf nodes to have their own classifiers. This is due to leafs being the last nodes in a tree, having no child nodes. To circumvent this shortcoming a solution is proposed that treats leaf node as if it had left child with the same assigned class as a parent, and a right child with assigned every existing class in the training dataset except for the class assigned to its sibling (left child node).

### 5.1.3. Classification rules

When an unknown, new pattern is presented to the Balanced Tree, it traverses a path from a root to a leaf node in order to be classified or rejected. This path strongly depends on classifiers in each node and their classification decision. As it was described earlier each node is assigned certain number of classes that it represents. The main task of each node's classifier is to decide if the provided pattern belongs to internal class '0' or '1'. In other words it tries to determine to which set of classes this unknown elements is most similar to. After decision is taken the patterns is sent further to the left child node in case it was classified as '0' or right one if classified as '1'. Each subsequent classifier is more precise and better clarifies pattern's class affiliation. After reaching leaf node the final classification test is made. The classifier in a leaf node is trained in an one-versus-all manner. If the unknown element is recognized as a member of a class assigned to this particular leaf, it is finally labelled as an element from that class. On the other hand if it is classified as a "rest" pattern, it gets rejected. The scheme for this approach can be seen on Figure 5.2. Rejection relies on the assumption that if the pattern traversed path all way down to the leaf node, while being sent to next nodes basing on increasingly strict classifiers' decisions, and ends up being recognized as a point from outside of most probable class (the one assigned to the leaf node), then it probably is not similar enough to any class from the training set.




Figure 5.2: Balanced Tree rejection scheme in leaf node

### 5.1.4. Implementation details

Creation of Balanced Tree structure starts from tree root and is done recursively. Each node, that is not a tree leaf, is assigned certain set of classes which is a subset of all classes in a tree (root node is assigned all). The next step involves clustering method dividing node's class set into two disjoint sets. This procedure is done on 'class central points' which are average points of all elements in each class. Clustering algorithm divides those points thus providing two new sets for both child nodes. After that node trains its classifier on data set consisting of two classes created

by taking all elements from training data for left and right child nodes' classes sets. The node-creation procedure is then applied for both node's children. The leaf creation algorithm is slightly different as it does not need usage of clustering. Classifier is trained on data set created from combining elements from training data that belongs to the same class the leaf node represents (those points' new class is labelled '0') and elements from every other class (which are labelled '1'). To ensure that both '0' and '1' classes have the same number of entries the '1' class set must be trimmed. This is done at its creation step by taking less elements from each class in order to have the same number (or nearly identical) of elements overall in the whole set, e.g. having training data set consisting of ten classes labelled from '0' to '9', with total of 10,000 elements, set '0' for leaf representing class '2' will have 1,000 entries of elements from class '2' taken from training data and set '1' will have 999 elements in total but will consist of elements from classes '0', '1', '3', '4', '5', '6', '7', '8', '9' taken from training data with 111 elements from each class.

## 5.2. Slanting Tree

### 5.2.1. Structure

The Slanting Tree structure differs greatly from Balanced Tree's one. The concept implemented in Slanting Tree assumes that the unknown pattern, that is sent for classification, should be iteratively compared against each class representatives. Only when it is similar enough to points from certain class, more precise tests are made that ensure its affiliation. Should the tests fail, the pattern continues its iteration over other classes as if wasn't ever supposed to belong to this class. Rejection occurs when every test fails.

Construction of the Slanting Tree is fairly simple, unlike the Balanced Tree. Each node represents exactly one class. Nodes are chained together in such manner that when traversing Slanting Tree from the root to the last node by choosing always the left child, exactly one node for each class in the training set is visited. Each of these nodes has also a right child that can be treated as a node between its parent and its parent's left child, that extends the path received when going from root to the last node by taking always the left node's successor. Each right child node represents the same class from the training data set as its parent does.

### 5.2.2. Classifiers creation

Each tree node in Slanting Tree has its own binary classifier, trained in a 'one-versus-rest' manner. Training is done on data containing two classes, where the first one consists of patterns

from the training set which belong to the same class as the node represents, and the second one is obtained by concatenating patterns of each class from the training set except for the class represented by the node. To prevent the situation in which node and its right child have the classifier trained on the same data (because both nodes represent the same class) certain changes to training patterns must be introduced. This ensures that every classifier in the Slanting Tree is unique and can be used during classification procedure.

### **5.2.3. Classification rules**

The classification starts from the root node, where the unknown pattern is tested by the first classifier and has its class affiliation checked. If the obtained result indicates that it isn't similar to the class represented by this node (gets classified as an element from the 'rest' class), the classification process is continued in the next node, which is the left child of the current one. If the opposite situation occurs, and the classifier accepts presented pattern as a representative of current node's class, the process is repeated in the right child node which uses more strict classifier. This is done to ensure that the unknown element, which is supposedly from the certain class, really belongs to it. If this test fails, the pattern is sent to the node as if the previous test also failed (it is sent to the left child of the current node's parent). In case of success the pattern gets successfully classified. When all tests fail (there is no more nodes to send pattern to) the element gets rejected.

### **5.2.4. Implementation details**

Creation of Slanting Tree is done recursively, starting from the root node. All classes that should be distinguishable by this tree structure are sorted by their labels and stored in an array object. This object is later used during node creation method to check what classes have already been covered by previous nodes. Every non-leaf node represents only one native class and has its binary classifier trained in 'one-vs-rest' manner, the same way the tree leafs' classifiers in Balanced Tree are (see 5.1.4). The next step involves creating left child node for the next native class in the array object that has not yet been used. In case of no classes left the function returns without creating new node. The last step consists of right child creation, which is a leaf node. Leaf nodes in a Slanting Tree represent the same native classes their parent node did, but their classifiers, although built using same 'one-vs-rest' approach, are trained on a different data sets in order to create more accurate results. Usually trained classifier does not achieve 100% accuracy even on a training test that was used during its creation. There are some samples from first class that get classified as elements from the second and vice versa. Such mistakes can help determine

what kind of corrections can be made to the classifier. For every non-leaf node, after its classifier training, there's set of elements from the first class that were correctly recognized (those are the elements from the class this particular node is representing) and set of elements from the second class that were mistakenly recognized as elements from the first class. Those two sets are used in this node's child leaf node's classifier creation. Of course before training those two sets must be the same size, ideally having the same number of elements as two sets used in parent's classifier training. For each missing element in either of sets the new object is generated by randomly selecting one element from this set and applying normal distribution (with standard deviation 1) to all of its features in a feature vector, thus getting new sample that can be added to the set. In case of having less than certain number of elements (implementation checks for 10 or less elements) in either of sets before new point generation algorithm takes place, those sets are filled with randomly selected points from parent node's classifier training sets.

### 5.3. Slanting Tree with ordered classes

#### 5.3.1. Description

The basic Slanting Tree structure assigns classes to its nodes based on an arbitrary, lexicographic order. This approach leaves its implementation vulnerable to situation in which label changes occur. Slanting Tree with ordered classes tries to circumvent this disadvantage by sorting classes without the need to know their labels, based on their spatial relations. Every set of points belonging to certain class in a training dataset can be transformed into one point, being the centre of that particular class. Next the distance to every other centre point in the training dataset is calculated for each class centre, and only the lowest value is saved. The new class order is based on those values, which are sorted in the descending order.

#### 5.3.2. Implementation details

Steps required to build Slanting Tree with ordered classes are mostly the same as in 5.2.4. Instead of creating nodes for classes using their lexicographic order, the ordering technique described in the previous Section is used. During computations only one point, called the class central point, for each class in the training dataset is used. Those are calculated the same way as in 5.1.1, by getting the average value of all training patterns from one class. Sorted class labels are used further during classifier tree creation process the same way as in original Slanting Tree.

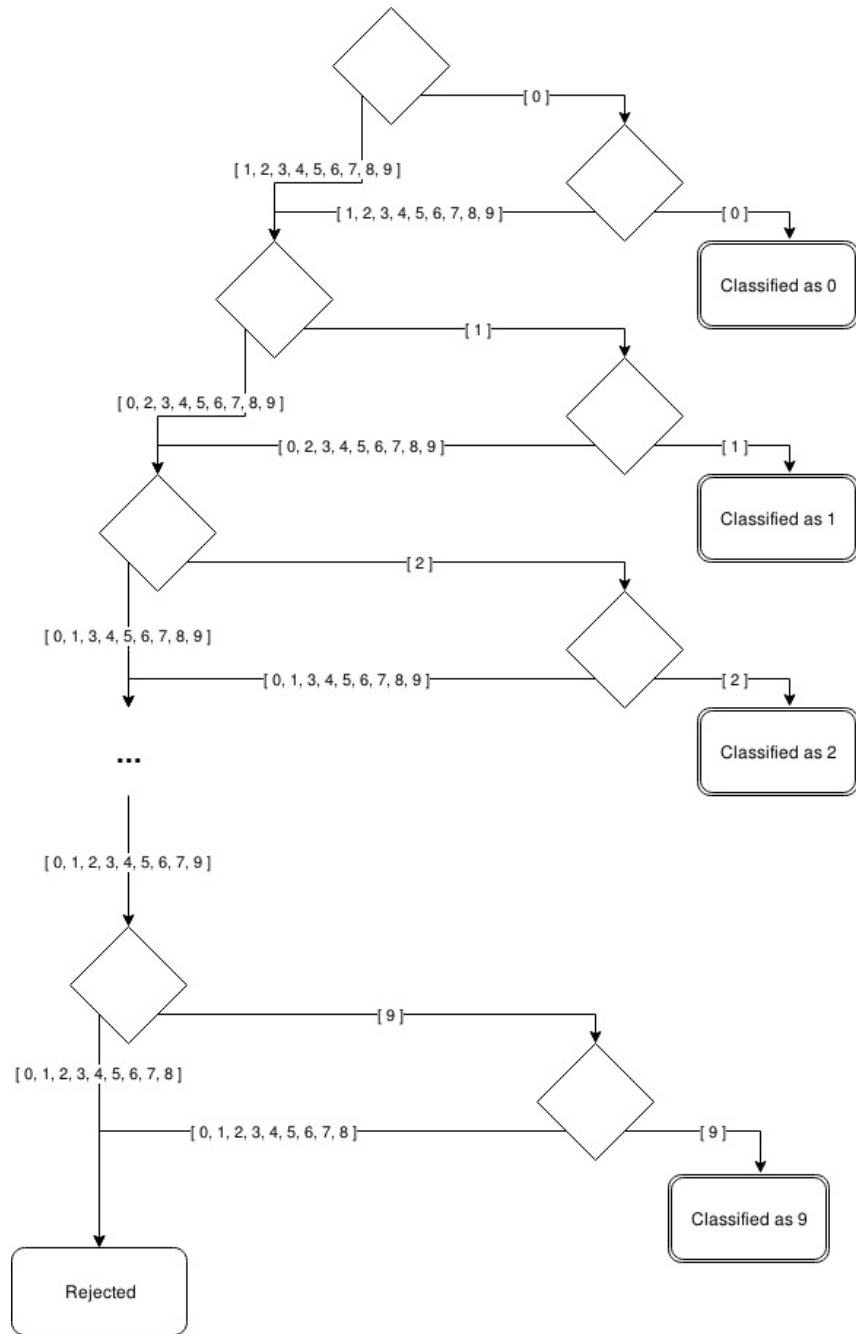


Figure 5.3: Bottom part of the Slanting Tree with nodes for classes 8 and 9 (nodes for classes from 0 - 7 not visible).

## 5.4. Slanting Tree 2

### 5.4.1. Description

Much like previously described Slanting Tree, this one has its nodes arranged in the same architecture. The difference lies in leaf nodes which, unlike the original Slanting Tree, are not using modified training data sets and use different classifier types instead (e.g. parent nodes using SVM classifier and their right children nodes using random forest). The idea behind this implementation relies on the assumption that various classifiers tend to wrongly classify different patterns, so when combining them rejection rate as well as classification rate should be vastly improved. Other than that there are no further changes and everything described in the Section 5.2 applies to Slanting Tree 2.

### 5.4.2. Implementation details

Creation procedure is mostly the same as in 5.2.4. The only differences are present in nodes creation method where instead of creating new training patterns for the right child node, different classifier type is trained on the same data from the parent node.

## 5.5. Results

Described in this chapter classifier trees were tested with various common classifiers: SVM, kNN and random forest, using different parameters. Over 500 tests were held. Results in form of quality measurements (see Chapter 3) for training, test and letters sets were gathered in form of two matrices with 12 rows and 10 columns, one for training and one for test data. Each row in the matrix corresponds to one of the quality evaluation measurements, and each column represents value of corresponding measurement scored by certain classifier tree using one of the common classifiers. See Table 5.1 for reference.

Every common classifier that was used by any of tree nodes was tested with different parameters. SVM had its C, gamma and kernel options adjusted (see Chapter 2 for every parameter explanation). Values were as follows

$$C : [1, 2, 4, 8, 16]$$

$$gamma : [2^{-1}, 2^{-2}, 2^{-3}, 2^{-4}]$$

$$kernel : [rbf, poly]$$



Table 5.1: Example empty result matrix

	Balanced Tree			Slanting Tree			Slanting Tree 2		
	kNN	SVM	RF	kNN	SVM	RF	kNN	SVM	RF
<b>Strict Acc.</b>									
<b>Fine Acc.</b>									
<b>Strict Native Sens.</b>									
<b>Acc.</b>									
<b>Native Prec.</b>									
<b>Native Sens.</b>									
<b>Native F-measure</b>									
<b>Foreign Prec.</b>									
<b>Foreign Sens.</b>									
<b>Foreign F-measure</b>									

Adjustments for kNN were made for only one parameter, using euclidean metrics

$$n\_neighbors : [3, 5, 7, 10]$$

Random forests also had modifications applied to one parameter

$$n\_estimators : [30, 50, 100, 150]$$

When evaluating results quality evaluation measurements were taken into account (see Chapter 3). In the next few subsections there is short summary for each classifier tree using different internal classifiers. Because the results obtained for Slanting Tree with ordered classes were almost exactly the same as for regular Slanting Tree (with arbitrary class order) the results tables do not include scores achieved for Slanting Tree with ordered classes to maintain their clarity.

Results gathered in Table 5.2 and Table 5.3 prove that combining commonly used classifiers by putting them in more complex structures does not affect overall classification capabilities. Of course the quality of determining patterns' affiliations relies mostly on the type of the classifier used, but is also affected by classifier's parameters and the tree structure.

Trees using SVM classifier yield better results when using radial basis function (rbf) kernel along with C parameter set to 16 and  $\gamma$  to 0.5. During calculations it was observed that the  $\gamma$  parameter didn't have as much impact on final results, unlike the C parameter which, when decreasing its value, lowered achieved scores. For Slanting Tree 2 the SVM classifier performed

Table 5.2: Measures values (described in Chapter 3) for classifier trees using various common classifiers on training data (described in Chapter 4)

	Balanced Tree			Slanting Tree			Slanting Tree 2		
	kNN	SVM	RF	kNN	SVM	RF	kNN	SVM	RF
<b>Strict Acc.</b>	20.67	44.92	41.47	22.76	29.70	50.29	38.34	41.51	40.59
<b>Fine Acc.</b>	92.66	99.33	100.00	90.61	91.79	92.43	94.08	95.41	95.94
<b>Strict Native Sens.</b>	92.64	99.09	100.00	90.00	91.73	92.43	93.06	95.26	95.79
<b>Acc.</b>	22.21	45.06	41.47	24.72	31.42	51.88	39.57	42.47	41.44
<b>Native Prec.</b>	21.23	27.60	26.38	21.70	23.41	30.35	25.62	26.69	26.35
<b>Native Sens.</b>	99.99	99.76	100.00	99.33	99.93	100.00	98.91	99.84	99.84
<b>Native F-measure</b>	35.02	43.23	41.74	35.62	37.93	46.57	40.70	42.13	41.69
<b>Foreign Prec.</b>	99.76	99.79	100.00	96.51	99.86	100.00	98.81	99.85	99.84
<b>Foreign Sens.</b>	1.58	30.55	25.94	4.92	13.25	39.11	23.82	27.25	25.94
<b>Foreign F-measure</b>	3.10	46.78	41.20	9.37	23.39	56.23	38.39	42.82	41.18

Table 5.3: Measures values (described in Chapter 3) for classifier trees using various common classifiers on test data (described in Chapter 4)

	Balanced Tree			Slanting Tree			Slanting Tree 2		
	kNN	SVM	RF	kNN	SVM	RF	kNN	SVM	RF
<b>Strict Acc.</b>	10.79	37.04	32.81	13.41	21.18	44.21	30.72	33.94	32.75
<b>Fine Acc.</b>	91.86	96.44	95.56	88.89	91.49	90.36	93.45	95.02	94.56
<b>Strict Native Sens.</b>	91.77	94.03	93.23	88.00	90.97	89.03	91.33	92.80	92.67
<b>Acc.</b>	11.62	37.39	33.26	14.53	22.05	45.18	31.37	34.44	33.30
<b>Native Prec.</b>	10.35	13.77	13.03	10.59	11.53	15.54	12.73	13.24	13.08
<b>Native Sens.</b>	99.90	97.50	97.57	99.00	99.43	98.53	97.73	97.67	98.00
<b>Native F-measure</b>	18.75	24.13	22.99	19.13	20.66	26.85	22.53	23.33	23.08
<b>Foreign Prec.</b>	99.28	99.08	98.94	97.74	99.52	99.58	98.93	99.04	99.13
<b>Foreign Sens.</b>	1.58	30.55	25.94	4.92	13.25	39.11	23.82	27.25	25.94
<b>Foreign F-measure</b>	3.10	46.70	41.11	9.38	23.38	56.16	38.40	42.74	41.12

best when paired up with Random Forest which may indicate that both of those classifiers tend to misclassifying different patterns (hence better rejection option rates).

Using Random Forest classifier yields different scores depending on which tree structure is used. Whereas Balanced Tree performs best when using Random Forest with 30 estimators, both Slanting Tree and Slanting Tree 2 get better results while utilizing classifier with 100 estimators. Interesting may be the fact that the best performing Slanting Tree 2 uses Random Forest combined with SVM classifiers with the exactly same parameters' values as when using SVM classifier backed up by Random Forest one. In both cases results are very similar which proves that those classifiers tend to cooperate well.

Unfortunately, among all classifiers tested, the kNN one performs the worst when used in all of the three trees. While this doesn't mean that the classifications rates were very low, in fact the differences between scores achieved by using kNN and SVM or Random Forest were negligible, rejection option was almost non-existent. The best results, achieved by Slanting Tree 2, were obtained when using Random Forest as a second classifier. All results for kNN classifier for each of classifier tree described in this chapter, presented in the tables, were achieved when using *n\_neighbors* parameter value of 10.

## 5.6. Summary

All of the classifier trees introduced in this chapter had good classification capabilities, very similar to the plain common classifiers they used. It is worth noting that not only did the classification rate stayed the same, but also rejection capabilities were introduced. Among all classifiers combinations tested it was the Slanting tree using random forests with 100 estimators that performed the best. Tables 5.2 and 5.3 show score achieved by this tree structure. Although being the best, classification rate achieved by this particular Slanting Tree may not be considered good, as it's lower than 50%. At best it could be seen as mediocre. Despite trying different classifiers and their parameters combinations no better solution could be found while using tree structures described in this chapter. The final conclusion can be made that the classifier trees introduced in this paper do not perform well enough to be used as a valid rejection mechanism. While still maintaining high classification rates those structures are slower than other popular classifiers which questions their usefulness.

## 6. Classifier Arrays

Another approach towards classification with rejection option problem involves chaining classifiers trained on certain, very specific data. The array of those classifiers serves as a voting mechanism where every new and unknown pattern is presented to each classifier in this array and the classification (or rejection) decision is made based on the overall achieved score. All classifiers inside the array are binary ones but can be divided into two groups:

- one-versus-all - those classifiers are trained on two sets, where the first one consists of training patterns from certain class, and the second one is made of all patterns from the training set except for those from this certain class
- one-versus-one - every classifier is trained on training patterns from two different classes

### 6.1. One-versus-all




Figure 6.1: “one-versus-all” rejection method. Unknown pattern passes through an array of specially prepared classifiers, one for each class. If each classifier says it is not native, it is rejected.

#### 6.1.1. Description

The “one-versus-all” method requires creating an array of binary classifiers. Training data set for each classifier in this method consists of two sets: the first one (denoted as “class\_i”) holding all training data for certain i-th native class, and the second one (denoted as “rest”) being the result of a subset sum operation performed on the rest of the classes except for the class used in

class\_i set. One problem with such approach is that as a result of the subset sum, class “rest” could contain significantly more samples than “class\_i”. In such case the “rest” set is undersampled.

### 6.1.2. Implementation details

The actual classification with rejection is performed by presenting the unknown pattern to each of the classifiers from the array. When any classifier recognizes this element as a native one (belonging to class\_i), then the pattern is treated as a recognized one, and it is assumed to be native. In a case when all classifiers reject a pattern (all binary classifiers say that it belongs to set “rest”), it is treated as a foreign pattern and it is rejected. It is worth noticing that there is a possibility that more than one classifier recognizes the pattern as a native element. In such case randomly chosen class label is assigned to this pattern. The scheme for this method is sketched in Figure 6.1.

## 6.2. One-versus-one




Figure 6.2: “one-versus-all” rejection method. Unknown pattern passes through an array of specially prepared classifiers, one for each class. If each classifier says it is not native, it is rejected.

### 6.2.1. Description

The “one-versus-one” method requires preparing an array of classifiers, but this time it consists of  $\binom{c}{2}$  classifiers, where  $c$  is the number of native classes. Each classifier is trained on data consisting of two sets: the first one (denoted as class\_i) holding all training data entries for i-th native class, and the second one (denoted as class\_o) holding all training data entries for some other class (not the same as class\_i). In the end, there is one classifier for each pair of classes: 1 vs. 2, 1 vs. 3, ..., 1 vs. c, ..., (c - 1) vs. c.

### 6.2.2. Implementation details

Classification with rejection mechanism is based on presenting unknown pattern to each classifier in the vector and remembering their answers (e.g. classifier constructed for 1 vs.  $c$  classes can classify the pattern as belonging to class 1 or class  $c$ ). In the end, those answers can be summarized and for each pattern a  $c$ -elements array with numbers saying how many times this pattern was classified as belonging to class 1, 2, 3,  $\dots$ ,  $c$  can be formed. The pattern is rejected when the difference between two biggest values in the result array is smaller than two. In such case, it is assumed that the classifiers were highly uncertain as to which class should this unknown element belong to. Otherwise, the pattern is classified as an element belonging to the class which had the biggest value in the result array. The general scheme for this method is presented in Figure 6.2.

### 6.3. One-versus-one modified

The modified “one-versus-one” method is based on the “one-versus-one” method discussed in 6.2. The difference between those two methods lies in a rejection mechanism. In this method an unknown pattern is treated as a foreign element if the biggest value in the result array is smaller than  $(c - 1)$ . What it actually means, is that there must be a certain class that has always been chosen by a classifier whenever it was possible.

### 6.4. Results

All the results presented in Tables 6.1 and 6.2 were obtained for kNN using `n_neighbors` parameter with value 10, SVM using rbf kernel, having  $C$  value of 8 and gamma 0.5 and random forest consisting of 100 estimators.

The results obtained when using kNN classifier aren’t very satisfying. All foreign elements were treated as a native ones which means that rejection option is not present in this approach. Both SVM and random forest performed better, but not as good as the classifier trees described in Chapter 5. Whereas 1 vs. all approach lacked definitive rejection option, the 1 vs. 1 was too strict in rejecting presented patterns. Although the last method, denoted as modified 1 vs. 1 approach brought balance to classification-rejection problem it didn’t manage to preserve high ratios for both options.

Table 6.1: Measures values (described in Chapter 3) for classifier hierarchy arrays using various common classifiers on training data (described in Chapter 4)

	1 vs. all			1 vs. 1			1 vs. 1 (2)		
	kNN	SVM	RF	kNN	SVM	RF	kNN	SVM	RF
<b>Strict Accuracy</b>	17.07	23.38	26.94	79.04	66.03	69.84	19.90	34.79	33.01
<b>Fine Accuracy</b>	81.32	91.23	91.34	-	99.70	100.00	94.94	98.50	100.00
<b>Strict Native Sensitivity</b>	81.32	91.17	91.34	0.00	9.58	9.33	94.94	98.46	100.00
<b>Accuracy</b>	20.99	25.22	28.75	79.04	66.04	69.84	20.96	35.10	33.01
<b>Native Precision</b>	20.97	21.88	22.73	-	11.82	14.92	20.96	24.41	23.83
<b>Native Sensitivity</b>	100.00	99.93	100.00	0.00	9.61	9.33	100.00	99.96	100.00
<b>Native F-measure</b>	34.66	35.90	37.04	-	10.60	11.48	34.66	39.23	38.49
<b>Foreign Precision</b>	100.00	99.65	100.00	79.04	77.16	78.13	-	99.94	100.00
<b>Foreign Sensitivity</b>	0.04	5.41	9.86	100.00	81.00	85.88	0.00	17.90	15.25
<b>Foreign F-measure</b>	0.08	10.26	17.95	88.29	79.04	81.82	-	30.36	26.46

Table 6.2: Measures values (described in Chapter 3) for classifier hierarchy arrays using various common classifiers on test data (described in Chapter 4)

	1 vs. all			1 vs. 1			1 vs. 1 (2)		
	kNN	SVM	RF	kNN	SVM	RF	kNN	SVM	RF
<b>Strict Accuracy</b>	8.28	14.09	17.88	89.78	73.72	78.11	9.47	25.93	23.40
<b>Fine Accuracy</b>	80.69	90.92	89.39	-	97.02	96.39	92.64	96.92	95.51
<b>Strict Native Sensitivity</b>	80.69	90.31	88.35	0.00	9.75	9.79	92.64	96.44	94.97
<b>Accuracy</b>	10.26	15.01	18.95	89.78	73.75	78.14	10.22	26.24	23.85
<b>Native Precision</b>	10.23	10.68	11.10	-	5.68	7.57	10.22	12.13	11.78
<b>Native Sensitivity</b>	100.00	99.33	98.83	0.00	10.05	10.15	100.00	99.50	99.43
<b>Native F-measure</b>	18.55	19.29	19.96	-	7.26	8.67	18.55	21.62	21.07
<b>Foreign Precision</b>	100.00	98.62	98.67	89.78	88.78	89.36	-	99.68	99.58
<b>Foreign Sensitivity</b>	0.04	5.41	9.86	100.00	81.00	85.88	0.00	17.90	15.25
<b>Foreign F-measure</b>	0.08	10.26	17.93	94.61	84.71	87.59	-	30.35	26.45

## 7. Identifier Arrays

The easiest and probably most intuitive way of dealing with classification task is to use patterns' spatial relations in order to determine their class memberships. This approach is used for example in kNN and SVM models, where point's affiliation is calculated based on its place in the features space. Every class in a dataset can be represented as a big "cloud" of points and usually, if the other clouds do not overlap each other, the more dense this cloud is, the easier the task of classification gets.

Each class' cloud can be enclosed in an arbitrary geometrical shape which can be used as an identifier. The difference between binary classifiers and identifiers is very subtle, yet important. Whereas binary classifiers can distinguish between patterns from two different classes, they must be trained on data consisting of elements from both classes. Identifiers accept (or one could say: identify) only those points that they were constructed on, and reject any outliers. Thus they require only one class to be provided during training process. One of the easiest and most intuitive model of identifier is minimum volume enclosing figure. Creating figure enclosing all elements that has the smallest volume possible ensures that the identification can be very strict, which in return helps to maintain high outlier rejection rate. As opposed to convex hull, which is the most accurate point set container with smallest volume and which is enclosed by linear hyperplanes, bounding figures are far less complex. In many cases, when there is a need for computing convex hull and testing inclusions of other points, an approximation of such hull can be used, which helps in reducing time needed for computations, since most of alternative methods have lower construction and inclusion-testing complexities. Among the most popular minimum volume enclosing figures there are: boxes, diamonds, simplexes and ellipsoids.

### 7.1. Minimum Volume Enclosing Ellipsoid

Minimum Volume Enclosing Ellipsoid (denoted as MVEE) problem is solved by several known algorithms that can be categorized as first-order, second-order interior-point or combination of the two. For small dimensions  $d$ , the MVEE problem can be solved in  $O(d^{O(d)}m)$  operations



using randomized or deterministic algorithms [8]. All the results presented in this chapter were obtained while using MVEE algorithm based on Khachiyan solution.

An ellipsoid in its centre form is given by the formula:

$$E = \{x \in \mathbb{R}^n | (x - c)^T A (x - c) \leq 1\}$$

where  $c \in \mathbb{R}^n$  is the centre of the ellipse  $E$  and  $A \in \mathbb{S}_{++}^n$  is a positive definite matrix. Points lying inside the ellipsoid satisfy

$$(x_i - c)^T A (x_i - c) \leq 1 + \varepsilon \quad (7.1)$$

where  $\varepsilon$  parameter defines the error margin in determining whether certain point belongs to ellipsoid. It can also be used “enlarge” the ellipsoid (by increasing  $\varepsilon$  value).

However, constructing minimal volume bounding ellipsoid is not a convex optimization problem. It turns out that the solution is not easily obtainable so the dual problem has to be found. For a more precise and in depth solution description see [8]. The main problem, when using ellipsoids as identifiers, lies in constructing them. Two main factors that decide about identification effectiveness are tolerance and acceptance parameters. Tolerance can be viewed as a threshold for ellipsoid construction accuracy. The lower the parameter is, the better minimal volume ellipsoid is created. On the other hand, even with a good training set, there is a risk of including native patterns that lie outside of the created ellipsoid. Acceptance parameter has been introduced to prevent such unwanted behaviour. It defines a threshold for point rejection for elements lying outside of the created figure.

## 7.2. Ellipsoids array classifier

The construction of a classifier constructed by using array of identifiers is pretty simple and straightforward. The array is filled with ellipsoids, one for each class in the training dataset. Every unknown pattern, sent for classification, moves through those ellipsoids and gets the information whether it lies inside the ellipsoid or not. In case of being part of identifier’s interior the value given by the equation (7.1) is taken into consideration, which tells about the position inside the ellipsoid (where value 0 means that the point is in the ellipsoid’s centre and 1 means that it lies on the surface). Finally, after all ellipsoids are checked, the one with the smallest equation (7.1) value (and for which the point lies inside it), designates the class to which this unknown pattern should be classified to. If no ellipsoid accepts the point, it is rejected and treated as a foreign element.




Figure 7.1: Array of ellipsoids that works as a classifier with rejection capabilities. Elements outside of existing ellipsoids are treated as a foreign patterns.

### 7.3. Results

The tests were performed and their results were stored in the same manner as in similar tests in previous chapters. The classification and rejection rates for identifier array using ellipsoids are very good, staying within 85% - 95% value.

Table 7.1: Results obtained for Ellipsoid arrays

	<b>Ellipsoid</b>
<b>Strict Accuracy</b>	88.28
<b>Fine Accuracy</b>	93.14
<b>Strict Native Sensitivity</b>	83.95
<b>Accuracy</b>	89.98
<b>Native Precision</b>	77.23
<b>Native Sensitivity</b>	90.13
<b>Native F-measure</b>	83.18
<b>Foreign Precision</b>	96.01
<b>Foreign Sensitivity</b>	89.93
<b>Foreign F-measure</b>	92.87

Although it may seem that the ellipsoids are superior to other classifiers presented and used in this paper it is worth noting that the tests were performed on only one dataset. Geometric classifiers that work on raw data, without introducing any changes to feature values given during training don't work well in situations when class elements overlap. SVM, random forest, etc. which should be used in such cases have the advantage over simple classifiers like kNN or arrays

of minimum volume enclosing figures, because they change the way features are interpreted. For example, SVM increases dimensionality by using kernel function to introduce such feature-space representation in which two different classes are easily separated by a hyperplane.

## A. An Appendix

## Bibliography

- [1] David Cournapeau. Scikit-learn website, 2007. URL <http://scikit-learn.org>.
- [2] Kuan Hoong. Kuan hoong blog, 2016. URL <https://kuanhoong.wordpress.com/2016/02/01/r-and-deep-learning-cnn-for-handwritten-digits-recognition/>.
- [3] Pedregos F. *Scikit-learn: Machine learning in Python*. Journal of Machine Learning Research, 2011.
- [4] Altman N. S. *An introduction to kernel and nearest-neighbor nonparametric regression*. The American Statistician 46 (3), 1992.
- [5] Vapnik V. Cortes, C. *Support-vector networks*. Machine Learning 20 (3), 1995.
- [6] L. Breiman. *Random Forests*. Machine Learning 45 (1), 2001.
- [7] Christopher J.C. Burges Yann LeCun, Corinna Cortes. The mnist database. URL <http://yann.lecun.com/exdb/mnist/>.
- [8] Yildirim E. A. Todd, M. J. *On Khachiyan's Algorithm for the Computation of Minimum Volume Enclosing Ellipsoids*. 2005. URL <http://people.orie.cornell.edu/miketodd/TYKhach.pdf>.